

Demystifying AI Session

22nd National Academy for IDEA Administrative Law Judges & Impartial Hearing Officers

Thursday, July 23, 2026

SPECIAL EDUCATION SOLUTIONS, LLC

(203) 557-6050

DMERCED@SPEDSOLUTIONS.COM

WWW.SPEDSOLUTIONS.COM

SESSION 1

I. INTRODUCTION

- A. This session focuses on understanding the growing prevalence, use, and application of artificial intelligence (AI) within special education. Specifically, the session will examine and provide recommended approaches for how ALJ and IHO leaders can better use and leverage AI as they go about their daily work processes, activities, and output.
- B. AI has moved from the margins to the mainstream of the legal profession. In 2025, for the first time, more lawyers reported using generative AI than not. Adoption among legal organizations nearly doubled in a single year, from 14% to 26%, according to Thomson Reuters, and 63% of small and mid-sized firms have now formally adopted it. At the same time, 81% of firm leaders still name reliability and risk as live concerns.
- C. In special education specifically, the shift is just as real. A 2025 Center for Democracy & Technology (CDT) brief found that 57% of teachers surveyed now use generative AI to help draft IEPs and 504 plans, a sharp jump from the year before. Parents, advocates, and attorneys are leaning on the same tools to draft due process complaints, analyze IEPs, and prepare for cross-examination. The filings, records, evaluations and various documents that reach you have increasingly passed through AI before they ever land on your desk.
- D. At the heart of this issue is an obvious tension. Adoption is climbing fast, and the cautions are real. If you are wary of AI, that is a good emotion to feel. That instinct is part of your job and well-founded. We have all seen the stories: lawyers sanctioned for citing cases that do not exist, briefs built on invented authorities, confident answers that turn out to be wrong. Healthy skepticism is the right starting point, and nothing today asks you to set it aside. What has changed are the tools. Today's models are markedly more accurate and easier to direct than the ones behind those headlines. Over just the past 6 months, tool capabilities (and accuracy) have grown dramatically. And now, there are concrete techniques for avoiding, or catching, the mistakes they may still make. The point of this session is not to sell you on AI. It is to show you how to get the benefits it offers, while also keeping the risks in check, and enabling you to maintain full control the entire way.
- E. To get you there, the session is built around five parts:

1. **Recognizing the AI Inflection Point:** See why AI is a genuine turning point, why it matters for your work, and why your skepticism is warranted.
2. **How AI Works:** Look under the hood at what AI is actually doing, so it feels less like a black box and more like a tool you can direct.
3. **Picking the Right AI Tool:** Compare the major platforms, see how they differ, and learn how to choose the right one for a given task.
4. **Putting AI to Work:** Apply AI to your daily responsibilities while protecting client confidentiality and personal information.
5. **Maximizing Your Productivity:** Rethink how the work gets done, not just how fast, so you can put this to use the day you are back at your desk.

By the time we close, you will know when AI is the right tool and when it is not, which platform fits your needs, how to write prompts that produce reliable results, and how to fold AI into your routine with confidence. You will also know how to avoid the failures that make headlines: fabricated citations, hallucinations, and the exposure of protected information. It's important you hold on to one principle through all of it. AI assists your judgment. It never replaces it. As an ALJ or IHO, you remain responsible for everything you review, decide, and sign, and nothing here changes that. The goal is not to hand your work to an AI machine. It is to give you a capable assistant that makes you faster and sharper while you stay firmly in control. Before we get to the five parts, it helps to step back and see the bigger picture: where AI fits in the long arc of technology, how it is already reshaping daily life, and why this moment is genuinely different.

II. RECOGNIZING THE AI INFLECTION POINT

- A. You're living through a once-in-a-generation shift. Over the last six centuries, only a handful of technologies have genuinely changed how people think, learn, and work. AI is the next one. When this session says AI, it means tools like ChatGPT, Claude, and Gemini, and we'll look under the hood at how they actually work in the coming sections.
 1. Monumental technologies impacting humans
 - a) **Printing Press** (1440): Mass access to knowledge
 - b) **Electricity** (1880s): Power on demand
 - c) **Combustion Engine** (1900s): Movement of people and goods
 - d) **Computers** (1940s to 1980s): Automation and scale
 - e) **Internet & Mobile** (1990s to 2000s): Instant connection and information
 - f) **AI** (2020s): Access to knowledge and reasoning ← **"You are here"**

Each wave wasn't really about the technology. It was about what the technology unlocked: who could reach knowledge, how goods and people moved, how fast we could compute, how instantly we could connect. Every one of them eventually faded into invisible infrastructure, the ordinary backdrop of how we live and work, and AI is on that same path. The only real question is whether you shape how it enters your work, or let it happen to you.

What AI redistributes is access to knowledge and reasoning, and it is doing so faster than any wave before it. The pace is the point. It is no longer a question of whether AI

reaches the legal field, but how fast. And, waiting on the sidelines is itself a decision with consequences. Skeptics of every prior wave were right in the short term and wrong in the long term. This is not a forecast. It is already happening, in client interactions and in hearing rooms. Your position, though, is unique. The attorneys and advocates before you can lean on AI as a drafting aid, but you are the neutral who has to weigh what they submit and sign a decision that shapes a child's education and a family's rights. So the honest concern is not whether AI is overhyped. It is whether a tool like this can be trusted in a setting where due process and a student's services are on the line. That concern is fair, and taking it seriously is exactly how you stay in control of the technology instead of the other way around.

B. AI is already in your world. You may not realize how much, but it is already woven throughout your field and the people around you. Here are the places you have probably seen it.

1. **School personnel**

- a) IEP drafting: PLAAFPs, goals, accommodations
- b) Prior Written Notices & meeting summaries
- c) Assessment data & progress narratives
- d) Document translation

2. **Parents, advocates & attorneys**

- a) Due process complaints & mediation requests
- b) IEP critiques & rebuttals
- c) Record summarization

3. **Legal research**

- a) AI inside Special Ed Connection, Harvey, Westlaw, Lexis+, CoCounsel
- b) Still not hallucination-free
- c) Special ed coverage lags

Whether or not you ever touch AI, the participants in your proceedings already do. 57% of teachers now use AI to help draft IEPs and 504 plans (CDT, Oct 2025), so the documents arriving on your desk have often passed through these tools first. Pro se complaints increasingly read like polished attorney work product, which blurs the line between represented and unrepresented parties. AI sits inside the very platforms you rely on and the fabricated-citation sanctions in the headlines show what happens when it is used without care. AI is already shaping the record before a case reaches you.

New York has addressed this directly. Under [22 NYCRR Part 161](#), “It is the policy of the Unified Court System that **the use by attorneys and parties of artificial intelligence tools in preparing papers submitted to a court should not be prohibited**, as long as such use is in accordance with the duties and responsibilities that apply to individuals who submit papers to a court. Since those duties and responsibilities already apply to all submissions, regardless of whether AI tools were used, attorneys and parties should **not be required, upon submitting papers, to disclose to the court that they have used AI** in the preparation of such papers.”

Part 161 speaks to attorneys and parties. Your position as the neutral is different, and courts and judicial ethics bodies have begun issuing AI guidance aimed at adjudicators. The themes line up with everything in this session: you may use AI as a tool, but you protect confidentiality, you verify what it produces, and you never delegate the decision itself.

So AI use, whether by you or by others, is not improper in itself. The real question is accuracy: making sure the documents you receive and the ones you produce are correct. Pair the legal experience you already have with the AI skills you'll build today, and you are well equipped for what is coming. Today is about your use of AI, using these tools well and safely. Let's take the mystery out of it. Once you see how it works, the potential becomes obvious.

III. HOW AI WORKS

- A. AI generates language. It doesn't look things up. Since search engines arrived in the late 1990s, looking something up meant typing a word or short phrase into Google and getting back a list of links. Over time people began typing whole questions, and today those questions have grown into detailed prompts with specific instructions. AI works differently: you tell it what you want, and it writes a response.
1. **SEARCH:** you enter a query, and it finds a page that already exists.
 2. **AI:** you enter a prompt, and it writes something entirely new.

At its heart, AI runs on Large Language Models (LLMs). These models are trained on enormous volumes of text, images, audio, and other content. An LLM predicts and reasons its way to an answer; it does not pull up indexed web pages that have been logged in a database. Ask AI the same question twice and you will often get two different answers, because the response is generated, not retrieved. Typically, Google hands you the same results every time, while AI generates an answer shaped to you and your question.

This generative nature is also why, without clear and specific direction, AI can sometimes invent things.

- B. Search vs. chat. The easiest way to see the difference firsthand is to send the same question to both an AI Chat and Google and watch what they provide as a response. Take a question every hearing officer knows: what is the difference between an IEP and a 504 plan? Type it into Google and you will likely get the familiar results page, a column of blue links that you need to open and read yourself (note, Google is in the midst of rolling out "AI Overviews" on select results pages). Put the same question to an AI chat and you get something different: a more detailed, synthesized, plain-language answer written for you, and a good AI tool will cite its sources so you can double check them.

That is the shift in one line: the search found pages, the AI wrote an answer. And it does not end there. You can reply in plain language, ask it to turn that answer into a

checklist a client could use, and it reworks what it just gave you. This back-and-forth is what surprises new users most. It is a conversation: you refine, redirect, and shape the output. You are directing the tool, not just querying it.

And typing is just the entry point. You can also talk to it out loud and have it answer back, or show it something live through your camera or screen. Want to have it review things or get additional context, just attach documents, spreadsheets, audio, video, or photos, and it will read, watch, or listen and then respond. It does not just answer questions in the abstract. It engages with you and with the material right in front of you, and for record-heavy work, that is where much of the value lives. Impressive. But you've also seen this technology stumble. Let's be honest about that.

- C. The common AI challenges and issues. You have read about these in the media, and you have probably hit a few yourself while experimenting. Here are the ones you are most likely to see.
1. **Hallucinations:** The model states something false with complete confidence. The fabricated case citations in the headlines are the obvious example, but it also shows up as wrong answers to simple questions, math errors, and output that is plain nonsense.
 2. **Knowledge Cutoffs:** A model is trained on data up to a fixed point in time, its training cutoff, so on its own it does not know about events, rules, or decisions that came after that date.
 3. **The "Averaging" Effect:** Because a model draws on patterns across all of its training data, it tends to hand you the safe, middle-of-the-road answer rather than the sharp or unconventional one.
 4. **Bad Prompting = Bad Results:** The output is only as good as the direction you give it. Vague input produces vague results.
 5. **Limited Memory:** By default, a model only works from what is in front of it. In a long exchange it can lose track of something you told it earlier, and it knows nothing about your case or your preferences unless you give it that context.
 6. **It Tends to Agree With You:** Models lean toward giving you the answer you seem to want. Ask one to support a position and it often will, which makes it a useful drafting aid but a poor substitute for an independent second opinion. The neutral judgment stays yours.

Keep in mind that AI does not store facts the way a database does. When it does not know something, and it has not been told to work from a specific source, it can fill the gap by inventing a reference or a fact. That is why, every time you use it, you have to read, review, and verify what it gives you before relying on it.

None of these problems is a reason to walk away from AI, because every one of them is improving, and quickly. The tools behind today's headlines are already noticeably better, and getting significantly more powerful as new models and capabilities are released. Let's now talk about these specific advances.

D. Dramatic AI improvements are occurring at hyperspeed. Model accuracy and capability are improving at a striking pace, and new models now ship from every major frontier provider on a roughly biweekly to monthly cadence.

1. **Model Advances:** The jump in raw capability over the last several months has been steep. Because these models are built on text and language patterns, software development was the first knowledge field to feel it, and the productivity gains there over the last two years have been dramatic. The same curve is now reaching other knowledge work, including legal, finance, and marketing.
2. **Context Capability:** One of the oldest limits was how much a model could hold in its “mind” at once: things such as earlier conversations, the documents you gave it, its own prior work. That working memory is the context window, and it has grown enormously over model iterations. Now, the leading frontier models leverage more than a million tokens. (A token is the model’s basic unit of text, roughly four characters, or about three-quarters of a word.) Providing proper context is a cornerstone of helping AI produce the best results.
3. **Live Research:** A model’s training data has a cutoff, often three to six months back and varying by model. As such, AI does not automatically know the latest guidelines, laws, regulations, or recent events. Never assume it is current on the information it is referencing. It’s important you tell it explicitly to search for the latest information, then vet what it finds.
4. **Source Constraints:** AI is far more reliable when it works from materials you give it. When you confine it to your own sources, using a Project, a referenced folder, or a tool like NotebookLM, hallucinations drop sharply. This is because the model is drawing on your documents or specific data, rather than its memory or whatever it happened to find on the web.
5. **Self-governance and Looping:** When you instruct it to, a model can now check its own work, fix what it finds, and validate the result before handing it back. Agents and subagents extend this further. An agent is an AI system that can take a goal and carry out the multi-step work to reach it on its own, deciding what actions to take, running them, and adjusting based on the results rather than waiting for a prompt at each step. Because agents operate with this autonomy, the same self-governance gets built into how they run, driving toward accurate results.

So the technology is ready. The real question now is a practical one: which tool do you actually pick? That is where we go next.

IV. PICKING THE RIGHT AI TOOL

A. The models are converging. For the most part, the major models, their capabilities, and their features are converging. They are starting to look a lot alike. As of mid-July 2026, the leading frontier models are:

- **ChatGPT** (OpenAI): GPT-5.6
- **Claude** (Anthropic): Opus 4.8 · Fable 5
- **Gemini** (Google): 3.1 Pro · 3.5 Flash

- **Grok** (xAI): Grok 4.3

The major frontier models are now broadly comparable and update on a roughly monthly basis. The providers are locked in an arms race over capability and performance, which is good for you, because it constantly raises the ceiling on what you can produce. Your own skill, and the value you get from these tools, will compound over time. And the principles you learn today will carry across all of them.

It is fair to ask how the big four actually differ. The honest answer is that the differences are now at the margins rather than the core: the interface and ecosystem, how each handles voice and images, what it connects to, and the feel of its writing. For most of the work you do, the more useful question is not which model is best, but which capable tool your organization provides and secures.

It should be noted, Grok has recently fallen behind its competition in performance, capabilities, feature set, and overall adoption across the enterprise community. As such, we will focus primarily on the other leading frontier models in the following.

Let's now look at how you can access and use these AI models.

- B. You have many tool choices. The landscape looks crowded, but it sorts into four broad types, and knowing what each type is for makes the rest of the picture click into place.

The four types are:

1. **General Chat** (web, desktop app, mobile app): This is the everyday workhorse, and for most people it is the whole experience of AI. In plain language you can have it answer questions, draft and summarize documents, explain a concept, translate, or think a problem through with you, and you sharpen the result simply by talking back to it. You can also hand it your own material, a document, a photo, even audio, and it will work directly from that. It runs on the web, your desktop, and your phone, which makes it the most flexible of the four and the easiest place to start.

Examples: ChatGPT (OpenAI) · Claude (Anthropic) · Gemini (Google)

2. **Workflow Assistants / Agents** (web, desktop app, mobile app): These go a step beyond a single answer. Instead of asking one question at a time, you hand the tool a whole job and it carries out the steps needed to finish it, working across your documents and even checking its own work along the way. You might point one at a folder of records and ask it to assemble a prehearing prep packet, and it returns a draft deliverable rather than just a reply. This is the “delegate” level of AI: the tool does the assembly, and you do the judging.

Examples: ChatGPT Codex (OpenAI) · Claude Cework (Anthropic) · Gemini Spark (Google) · OpenClaw · Hermes · Perplexity · Genspark · Manus

3. **Built into Office Tools:** This is the same kind of AI capability, but it lives inside the applications you already use, Word, Outlook, Excel, Gmail, Docs, and the browser, instead of a separate window. It works on your real documents and email in place, so there is no copying back and forth between apps. And because it runs inside your organization's licensed environment, its privacy and security protections are already built in, which matters enormously for confidential work. It also tends to be the lowest-friction option, since you may already have it or be a single subscription tier away.

Examples: Microsoft Copilot · Google Workspace

4. **Legal-specific:** These are purpose-built for legal work and grounded in curated databases of case law, statutes, and secondary sources. Because they draw on real legal material rather than a general model's memory, they can cite actual cases and authorities and reason from the law itself, which is precisely what guards against the fabricated-citation problem that has burned attorneys in the headlines. They also understand legal tasks out of the box, from research to citation-checking to document review. For anything that touches case law this is the right home, with the trade-off that these tools tend to be the most expensive of the four.

Examples: Harvey · Westlaw AI / CoCounsel (Thomson Reuters) · Lexis+ AI (LexisNexis) · Claude for Legal (Anthropic)

Whichever tool you choose, the next decision is the same: free or paid. The answer is easy.

- C. What to consider. What to use. Nearly all of these tools are sold in a few standard ways:

1. Charging for Usage

- a) **Free:** Limited access, usually to less powerful models and fewer features.
- b) **Monthly or Annual Subscription:** A set allowance of usage, with access to the latest models and the full feature set.
- c) **Usage Volume:** Pay as you go by the token, much like a utility meter.

2. User Account Types / Plans

- a) **Consumer / Small Business:** Built for individuals and small offices with simpler IT and security needs. The catch: consumer plans typically train on what you type, retain your data, and may allow human review.
- b) **Enterprise / Teams:** Built for individuals or organizations with real security, privacy, and data-management needs, and centralized IT control. The key difference: enterprise plans do not train on your data, and that protection is written into a signed agreement.

Free AI tools and most consumer-oriented plans should not be used in your role as an ALJ or IHO. Given the stakes of your work and your duty to protect client privacy and confidentiality, the security and data-training pitfalls of the consumer-grade tools are not worth the risk.

So what should you actually use? If you work within an organization or a government department, you most likely already have enterprise-grade tools at hand, such as the Microsoft Office Suite or Google Workspace. Your IT department controls access and manages security under established policies. If you want to use a tool they have not provided, consult IT first about integrating it securely.

If you buy your own software, the safest and most efficient path is to use the AI built into the office tools you already run, such as Microsoft 365 with Copilot or Google Workspace. These come with the proper privacy and security protections by default. You may find the AI features are already switched on, or that a small subscription upgrade unlocks the more advanced ones.

You might be wondering which models power those two office suites. Microsoft's Copilot takes a multi-model approach, drawing on OpenAI's GPT models, Anthropic's Claude, and Microsoft's own in-house models. At present, Microsoft products default to the models from ChatGPT. In the coming month (August 2026), Microsoft is planning to roll out the ability for users to select their model of choice (ChatGPT or Claude). Google's Workspace runs on its own Gemini models, alongside specialized media and reasoning models tuned for enterprise work.

For you, a subscription plan is the only sensible choice, and it is a genuine value. Instead of paying per token, which adds up fast, a subscription lets the providers effectively subsidize your usage, giving you something like five to twenty times the volume for the same spend.

That is the first half: what AI is, how it works, and what to buy. The second half is about putting it to work.

SESSION 2

V. INTRODUCTION

- A. Let's get to work. You now know what AI is, how it works, and how to pick a tool. Now we'll focus on how you can apply and use it. Specifically, we'll dive into three areas: 1) Putting AI to Work, the specific places it fits into your day; 2) Maximizing Your Productivity, rethinking how the work gets done instead of just doing the old tasks faster; and 3) Getting Started on Monday, a single first step you can take the moment you are back at your desk. By the end, you will know how to put AI to use right away, leverage its capabilities with more confidence, and produce better results with less effort.

VI. PUTTING AI TO WORK

- A. Where can AI help you? Before we get into that, one thing has to come first: protecting client confidentiality and personally identifiable information (PII) is critical. The next

section covers the specific approaches and tools for doing exactly that, so keep it in mind as you read what follows.

Given how fast AI's capabilities are expanding, and the way your daily work actually gets done, there are four functional areas where it can help you the most.

1. **Document Review & Summarization**
 - a) Organizing and summarizing exhibits, reports, and records
 - b) Building chronologies and hearing/meeting prep materials
2. **Analysis & Accuracy Checking**
 - a) Reviewing party submissions and spotting gaps
 - b) Verifying dates, citations, and facts against the source
 - c) Thinking / brainstorming partner
3. **Drafting & Communications**
 - a) First drafts of decisions and orders, in your structure
 - b) Scheduling orders, correspondence, and templated responses
 - c) Plain-language versions for parents and parties
4. **Docket & Calendar Management**
 - a) Calendar summaries and day preparation
 - b) Scheduled routines: news digests, weekly deadline summaries

Before our session, consider whether these four areas capture your daily work. Are there other activities or work products you handle regularly that you would add? What specific pain points do you run into? Where else could AI make your work more efficient?

This is a comprehensive list, and we'll return to these four areas later in this session to show what putting each one into practice actually looks like.

But none of that starts until we handle one thing. Before we touch a single case file, we protect what matters.

- B. The first step, protecting what's important. Protecting PII is central to your role. Every day you receive reports, evaluations, interviews, transcripts, video, audio, and documents packed with information that has to stay confidential: student names, school and district names, and anything else that could identify a student, their characteristics, or their disability.

Recall how AI models train on the material they take in. Many also keep and train on what you give them directly: your chats, the documents you upload, the audio and video you share. On free and consumer plans, that is usually the default, not the exception. That is why you should never use a free or consumer-grade AI account to discuss your cases, upload confidential documents, or enter anything not meant for the public.

As we covered in Session 1, the answer is to work inside enterprise-grade tools your organization has procured under a signed agreement, such as Microsoft 365 Copilot, Google Workspace, or the Team, Enterprise, or Edu tiers of ChatGPT or Claude. On those plans your data is excluded from training by contract, not by a toggle you have to remember to flip.

With the right tool, the protection comes from the boundary, not from scrubbing the documents. Inside a properly licensed enterprise environment, the AI works within your organization's existing data boundary; your prompts, the files you provide, and the responses it generates stay inside that boundary and are not used to train the models or shared with anyone else. That is what lets organizations run these tools on complete files without redacting first. Before you upload or expose any client PII, confirm with your IT department or the vendor that the data-handling, retention, and no-training protections are actually in place for the plan you are on.

Protecting what goes in is only half the job. The other half is controlling what comes back out. No setting or toggle reliably stops an AI tool from surfacing PII in what it generates, so this part is on you.

Your options, from weakest to strongest:

1. **In-chat instruction:** Tell the model, in the conversation, to leave out names and identifying details (or, replace these identifiers with more generic terms). It works most of the time, but the model can slip, so you read every output before using it.
2. **Standing instruction:** Put the no-PII rule into a saved workspace's instructions so it applies automatically to every chat there, instead of retyping it each time. More durable, same reliability ceiling: you still verify. Again, you can instruct the AI to replace identifiers with more generic terms.
3. **Redact before upload:** Strip the identifying details out of the document before it ever reaches the tool. It is the most work, and the only method that actually guarantees the result.

Whatever tool you land on, run the same short checklist before you feed it real case material:

- **Confirm training is off.** On enterprise plans the agreement handles this; on anything else, check the setting yourself, because it is often on by default.
- **Scope its access.** Point the tool at specific folders only, never your whole computer or drive.
- **Use governed storage.** Keep files inside your organization's environment (Copilot and Office, Workspace) or another designated secure location.
- **Organize by matter.** Use one Project per case so context stays clean and case information or details don't bleed into each other.
- **Read every output.** No tool removes this step, and it is the one that catches what the others miss.

With the guardrails in place, we can turn to the skill that shapes everything the AI gives back: telling it clearly what you want.

- C. Prompting is key to success. The quality of what you get back depends almost entirely on the quality of the direction you give. Ask the same tool the same question two different ways, and you can get a vague paragraph or exactly the analysis you needed. The encouraging part is that good prompting is not a talent or a trick. It follows a simple, repeatable pattern, and once you have it, it works on every tool.

The optimal AI prompting method is the **R-C-T-F-C framework**. It enables you to create a strong, comprehensive prompt by breaking it into five distinctive parts:

1. **Role:** Who you want the AI to be. For example, cast it as an experienced New York special education attorney and it takes on the right vocabulary, assumptions, and tone.
2. **Context:** The background AI needs to be useful. What you are working on, who it is for, and the facts or limits that matter.
3. **Task:** The specific thing you want done, stated plainly. One clear instruction beats a general request.
4. **Format:** How the answer should be structured. A bulleted list, a table, a draft order, document structure, headings and any other expectations you have.
5. **Constraints:** The boundaries it must respect. Length, tone, what to leave out, which sources to rely on or avoid, and anything else you explicitly do not want.

You will not need all five every time, and they do not have to come in order. But the more of them you provide, the closer the first response lands to what you actually want.

To provide you with a tangible example, here is what the five parts look like for a task you might actually run:

- **Role:** “You are a New York attorney with 15 years of experience in special education law.”
- **Context:** “I am an impartial hearing officer preparing for _____. Use the following documents as your source materials: _____”
- **Task:** “_____.”
- **Format:** “Present the answer as a bulleted list with bold headings.”
- **Constraints:** “Keep the response under 500 words and avoid technical jargon.”

Strung together, one after another, those five lines are a single, complete prompt. You will build exactly this kind of prompt at the end of the session, in the starter prompt you can put to use on Monday.

Good direction is the core skill. On top of it, the platforms hand you a set of tools that multiply what a good prompt can do.

- D. Work smart, work efficiently. A good prompt gets you a good answer once. The real efficiency comes from not having to rebuild that prompt, or re-supply your context, every single time. Every major platform now gives you a set of reusable tools that let you set something up once and then use it again and again.

These tools sit a step beyond everyday chat, and we are introducing them on purpose. This is where the entire field is heading. Over the coming year, these capabilities will move from optional extras to the center of how efficient professionals work. The encouraging part is that none of it requires technical skill. The platforms are making these simpler by the month, and most take only a few clicks to set up. You do not have to adopt them all at once, or even soon. For now, the goal is just to recognize them by name, know that they exist and potentially reach for them if it will make your job easier.

Here are the common platform (reusable) tools and what each one provides to you:

1. **Skills:** A saved set of instructions that teaches the AI to handle a recurring task your way, so you never re-explain it. You might build one that knows your decision format or your standard scheduling-order language, then simply call it up when you need it. They are written in plain text, which makes them easy to read, adjust, and even create yourself.
2. **Plugins:** Ready-made bundles you install to extend what the AI can do, usually packaging several skills, commands, and connections into one. Think of them as pre-built toolkits: instead of assembling everything yourself, you add one that already does the job. Many are built and shared by others, so you benefit from setup work you never had to do.
3. **Projects:** A dedicated workspace for a single matter, where the relevant documents and standing instructions live in one place. Because the AI always works from that fixed set of materials, its answers stay consistent and grounded, and, as covered in the PII section, keeping one Project per case stops matters from bleeding into each other.
4. **Connectors / MCPs:** Secure links that let the AI work directly with your other systems, such as your calendar, email, document storage, or a legal research source. Rather than copying and pasting between apps, the AI reads and acts where your information already lives. Access is permission-based: it can reach only what you can, and only what you choose to connect. (MCP, the Model Context Protocol, is the shared standard these links are built on.)
5. **Routines / Crons:** Scheduled tasks that run automatically on a timer, with no one at the keyboard. A routine might deliver a morning digest of new special education decisions, or drop a summary of the week's deadlines in your inbox before Monday. You set it up once, and it runs on its own.
6. **Browser Extensions / AI-specific Browsers (e.g. ChatGPT Atlas, etc.):** An AI assistant that lives in your web browser (extensions appear as a sidebar), able to see the page you are on and help with it directly, whether that is summarizing a long opinion, pulling the key points from a record, or filling out a form. It rides along with your normal browsing as a ready helper.

You do not need all six, and certainly not on day one. The thread running through every one of them is the same: set it up once, then stop repeating yourself. For most hearing officers, Projects and Routines are the natural place to start, a Project to keep each case organized and a routine to handle the recurring reading. And when you add Connectors, use only the governed ones your organization approves, for the same confidentiality reasons we covered earlier.

These tools make the work faster. But the bigger opportunity is not doing the old work faster. It is rethinking how the work gets done in the first place.

VII. MAXIMIZING YOUR PRODUCTIVITY

- A. Don't just speed up tasks. Rethink the work. Most people aim AI at the tasks they already do and get them done a little faster. The larger opportunity is to rethink how the work gets done at all. In Session 1, we discussed how software development showed the pattern: first the models answered coding questions, then they assisted with the work, then they took on whole jobs, and eventually they ran for stretches checking their own output.

Your work follows the same staircase, one step at a time:

1. **Step 1, ASK** · "What does ____ mean in the context of ____?" (*a question, an answer*)
2. **Step 2, ASSIST** · "Summarize this de-identified evaluation report" (*a task, a draft*)
3. **Step 3, DELEGATE** · "Build my prehearing prep packet from these documents" (*a workflow, a deliverable*)
4. **Step 4, ORCHESTRATE** · "Draft, then check every date and citation against the source, then flag gaps" (*a process that reviews itself*)

Most users stop at Step 2, and Step 2 alone is genuinely useful. But the gains that compound live at Steps 3 and 4.

Step 3 is where the tools from the previous section come together: a Project for the matter, a Skill that knows your format, and your own documents combine into a repeatable workflow rather than a one-off chat. Step 4 is the self-governance idea from Session 1, where you tell the model to check its own output against the source before it reaches you. You still verify; it just hands you cleaner work.

None of this changes the one rule: you review, you decide, you sign. The staircase changes how much assembly work happens before the judging begins. It never changes who judges.

You do not have to start at the top. Find the step where your work sits today, get comfortable there, and then climb one. Most of the gain comes from a single move up: from asking to assisting, or from assisting to delegating. The goal is not to reach Step 4 quickly, but to keep climbing deliberately.

That is the concept. The next question is where it applies most in your actual week.

- B. Where you can leverage AI the most. This is where AI can make a sizable impact on how you approach your work, spend your time, and get more accomplished in a given week.

You'll note, these are the same four areas we mapped out earlier, now seen from the other side. Not what AI can do in each, but rather the practical difference AI makes, across your entire spectrum of work activities.

1. **Document Review & Summarization.** Today, you spend hours assembling exhibits, summarizing reports, and building chronologies by hand. With AI, summaries and chronologies are drafted from your own documents; the reading-everything-twice burden drops, while the verifying burden stays with you.
2. **Analysis & Accuracy Checking.** Today, you re-read the record by hand to check dates, citations, and gaps. With AI, it cross-checks every date and citation against the source and flags what is missing, and you judge what it finds. This is the orchestrate step from the staircase, in practice.
3. **Drafting & Communications.** Today, every decision starts from a blank page, and scheduling orders and letters are written one at a time. With AI, you get a structured first draft in your own format, and routine correspondence from templates in minutes; you edit, decide, and sign. A rough first draft is faster to fix than a blank page is to fill, so if the rest feels like too much, start with correspondence: the most boring example and the easiest win.
4. **Docket & Calendar Management.** Today, you manually scan for new decisions and track deadlines. With AI, a morning news digest and a weekly deadline summary are waiting in your inbox, set up once as a routine.

What does not move is the heart of the job: findings of fact, credibility determinations, and the final legal language. Those were never assembly work; they are the work. AI buys back your time for exactly that.

So where do you actually start? With exactly one thing.

VIII. GETTING STARTED ON MONDAY

- A. A toolkit for AI success. Let's set you up for a strong first day. Here is a simple, step-by-step way to hit the ground running on Monday:
1. **Set up right:** Confirm you have the right productivity tools in place, with AI working alongside them (Microsoft or Google, as we recommended). Check with your IT staff if needed about storage, privacy, and data handling.
 2. **Create one Project:** The best way to learn is to dive in. Set up one Project for a single matter, then load its source material and context.

3. **Run the starter prompt:** Take one de-identified document and run the starter prompt below on it. Pick something that normally eats your time but has a clear, well-defined outcome.
4. **Verify everything against the source:** Read the output against the original, and note what it got right and what it got wrong.

The Starter Prompt: "You are an experienced legal assistant supporting a special education hearing officer. I will paste one de-identified document from a due process matter. Summarize the document, identify the issues it presents, and list every date or deadline it references. Format your answer with bold headings: Summary, Issues Presented, Dates & Deadlines, Open Questions. Use only the document provided; do not cite case law or outside sources; do not include PII, replace with generic terms. If something is not stated in the document, write 'not stated.' Keep it under 750 words."

Not ten things. One thing. One document, one prompt, one verification pass. Twenty minutes on Monday.

Notice that the prompt is simply the R-C-T-F-C framework from the prompting section, filled in: a Role (legal assistant), Context (a de-identified due process document), Task (summarize, list the issues and the dates), Format (bold headings), and Constraints (use only the document, no outside case law, write "not stated" when something is missing, stay under 750 words). That constraints line is doing the safety work, the same source-grounding idea from Session 1 written directly into the prompt.

And the verification step is the real lesson. Reading the output against the source teaches you the tool's strengths and its failure modes on your own material, at zero risk.

You are ready to put AI to work on Monday. That leaves just one thing: the rule to carry through all of it.

- B. The one rule to remember. AI is a powerful tool. It can help you learn, save you time, and make you remarkably more efficient. But there is one thing to keep in mind. *You are responsible for everything you review, create, and submit. AI assists you; you are ultimately accountable.*

We opened with a promise, and you now have it. You know:

- ✓ When to use AI
- ✓ When not to
- ✓ How to verify everything it produces

Across the two sessions, you saw why AI is a genuine turning point, how it actually works, how to choose the right tool, how to put it to work while protecting what is confidential, and how to rethink the work itself.

Everything else was technique. This is the principle. Hold to it, and AI becomes exactly what it should be in your work: a capable assistant that makes you faster and sharper, while the judgment, and the decision, remain yours.